

7, 425

LAB 3

grep - copy

sed - search and replace

*regular expressions

*learn regular expressions table for test week 6

*teaching notes 4.1

* is a wild card, you can use it as grep " " *.txt -> meaning search in any file in the given directory that ends with .txt

GREP

-> case sensitive

-> use -i for it to be case insensitive

-match all lines that contain ana

grep "ana" input.txt

ana

1234ana

stii unde e anabel

afbhid kna anastasia acbdi

grep -i "Ana" input.txt

ana

1234ana

stii unde e anabel

afbhid kna anastasia acbdi

SED

-> just outputs the result, without replacing the original file

-> s/ is for substitute

-> /d if like a filter, deletes all files containing the given expression

sed "s/ana/maria/" input.txt

maria

1234maria

stii unde e mariabel

afbhid kna mariastasia acbdi

nothing

12345

sed "/ana/d" input.txt

nothing

12345

sed 'y/aei/@3!/' input.txt

@n@

1234@n@

st!! und3 3 @n@b3l

@fbh!k kn@ @n@st@s!@ @cbd!

```
noth!ng  
12345
```

LAB 4

[abc] – takes them with an OR between them → will return a line that contains any of these characters
[a-z0-9] *matches ONE SINGLE character
+ → one or more times
* → zero or more times
 → can be looked at like a while loop
? → zero or more times, binary → yes or no
[a-z]{3} → the sequence has to be matched at least 3 times
^ → beginning of the line or not (^ana or [^a-z])

grep -E '/^ana/' n.txt → searches for ana at the beginning of the line

grep -E '/^.{3}\$/' n.txt → any line that has exactly 3 characters
 → '/^...\$' does exactly the same

cat /etc/passwd | wc -l → takes the output and counts the lines

grep -E "^[a-z]{2}[02468]+:" /etc/passwd
grep -E "^.{2}[02468]+:" /etc/passwd
 → find usernames that contain only even numbers

grep -E "^([:]*[aeiou][:]*{2}:" /etc/passwd
 → username has at least 2 vowels

grep -E "^([:]+){4}[:]*\<ION\>" /etc/passwd
 → find the students with the name ION

LAB 5

SED

> sed -E 's/ _____ / _____ /'
 ^ search ^ replace
> sed -E 's/ [A-Z] / - /' → replace the first occurrence on each line
> sed -E 's / [A-Z] / - / g' → the flag "g" stands for global, it replaces every occurrence
> sed -E 's / [A-Z] / - / gi' → replace everything, case insensitive
 → using -i replaces the content of a file, not just prints it

* replace vowels with *
> sed -E 's/[aeiou]/*/' file.txt → replaces just the first occurrence
> sed -E 's/[aeiou]/*g' file.txt → replaces all of them globally
> sed -E 's/[aeiou]/*gi' file.txt → replaces everything case insensitively

* add before each vowel !
 → grouping 'sed/([aeioy])()()/*\1\2/'
 > sed -E 's/([aeiou])/*\1/gi' file.txt → adds before every vowel
*

```
* swap uppercase-lowercase groups
    > sed -E 's/([A-Z])([a-z])/\2\1/g' file.txt *WITHOUT i

* double each vowel
    > sed -E 's/([aeiou])/\\1\\1/gi' file.txt

operator y/
    -> y/aeiou/@3!0w/
    -> mapping? not string replacement, but rather a correspondence

delete
    > sed -E /Ana/d file.txt -> deletes all lines that contain Ana
```

AWK

```
-> basically a mini programming language
-> print $smt prints a token that's split by a given delimitator

> awk '{print $2}' file.txt
    -> if that word doesn't exists, it prints a new line
    -> every line gets applied this print $2

* telling awk to define tokens by my given delimitator
> awk -F: '{print $2}' file.txt
> awk '{print NR, $0}' file.txt -> NR is the num of the line
> awk '{print NF, $0}' file.txt -> NF is the number of fileds
> awk '{print "start"} {print NF} {print "end"}' file.txt -> between
  '' we can have as many bodies as we want
> awk 'BEGIN{print "start"} {print NF} END{print "end"}' file.txt ->
BEGIN prints it just at the start and END prints it once at the end
```

You can have a __.awk file to format it better

```
> awk -f try.awk file.txt
1 BEGIN{
2     print "Welcome this is the beginning of the file"
3 }
4
5 NR%2==0{print $0}
6
7 END{
8     print "Okay byeee"
9 }
    -> print all lines that are even

1 BEGIN{
2     print "Welcome this is the beginning of the file"
3 }
4
5 NR%2==0&&length($2)>5{print $0}
6
7 END{
8     print "Okay byeee"
9 }
```

```

-> print all even lines for which the second token is > 5

* write an awk script that will count how many users are 914 and print it
1 BEGIN{
2     cnt=0
3 }
4 $5 ~ / 914 / {
5     cnt++
6     print $5
7 }
8
9 END{
10    print cnt
11 }

* piping, get the last name of people in 914 and replace vowels with *
> awk -F: -f ex.awk /etc/passwd | awk '{print $5}' | sed -E
's/[aeiou]/*/gi'

```

LAB 6

we have to be super precies with the interpreter

* for permission (user, group or other), use chmod +x filename or chmod 755 filename

```

#!/bin/bash -> this is the interpreter <-> /bin/bash a.sh
$1 -> first argument in the command line
$@ -> all arguments in the command line
$# -> number of arguments in the command line

```

tabs are important in bash, just like in python

\$ works like a pointer in bash

B=1 would add a value to a variable, while \$B would get the value of the variable

* if you want to use a variable in the command line, use \$variable

```

a.sh
#!/bin/bash
echo "Hello World, $1!"

for A in $@; do
    echo "$A"
done

```

LAB 7

for most of the requirements, you have to think that there must be a shell command that can do it (sed / grep / awk)

you can see grep as a function that provides filtering

```

for A in $@; do // A is iterator, $@ are all the arguments
    // TABS
    // check if they are files or directories

```

```

        if test -f $A; then // -f is for file -> this is the "shell way of
doing it"
            echo "$A is a file"
        elif [-d $A]; then // -d is for directory -> more like a switch
case in other languages
            echo "$A is a directory"
        else
            echo "$A is not valid"
        fi
done

```

if you want to verify lower branches, you need to use the full path of
the file

5.6.2.2

```

for A in $@; do
    if test -f $A; then
        echo "$A is a file"
    elif [-d $A]; then
        echo "$A is a directory"
    elif echo $A | grep -q -E '^[0-9]+$'; then
        echo "$A is a number"
    fi
done

```

5.5.2.1

count all lines from all files

```

for A in `find . -type f`; do
    // when you want to associate a value do a variable don't use
    spaces and "capture" the result with ``
    N=`grep -E '^[0-9]+$' $A | wc -l` // wc -l counts the lines
    S=`expr $S+$N` // expr is a shell command that can do operations
done

find all files in a directory that have a size greater than a given
value
the 5th field of ls -l is the size of the file
the classical operators do not work for comparing numbers in shell!!
    -gt, -lt, -eq, -ne, -le, -ge
D=$1
S=$2
for A in `find $D -type f`; do
    N=`ls -l $A | awk '{print $5}'``
    if test $N -gt $S; then
        echo $A
    fi
done

```

WHILE
read from keyboard until the user types "stop"

5.7.1.3

== verifies 2 strings

```

1#!/bin/bash
2
3 while true; do
4     read X
5     if [ "$X"=="stop" ]; then // "" are important, since they CAST
the variable to a string
6         break
7     fi
8 done

```

*script that receives a directory and checks for any changes into it
change can be at the file level or at the directory level, content or details

	content	details
file	cat f	ls -l f
directory	ls -l d	ls -l -d d

-> build hash of all files and directories

-> hash is one way, you cannot go back

Checking if the directory changed in any way

```
#!/bin/bash
```

```

D=$1
if [ -z "$D" ]; then
    echo "ERROR: No directory provided for monitoring" >&2
    exit 1
fi
if [ ! -d "$D" ]; then
    echo "ERROR: Directory $D does not exist" >&2
    exit 1
fi
STATE=""
while true; do
    S=""
    for P in `find $D`; do
        if [ -f $P ]; then
            LS='ls -l $P | sha1sum'
            CONTENT='sha1sum $P'
        elif [ -d $P ]; then
            LS='ls -l -d $P | sha1sum'
            CONTENT='ls -l $P | sha1sum'
        fi
        S="$S\n$LS $CONTENT"
    done
    if [ -n "$STATE" ] && [ "$S" != "$STATE" ]; then
        echo "Directory state changed"
    fi
    STATE=$S
    sleep 1
done

```

redirection errors

LAB 9

1.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

// create n child processes using fork and print pid of each child
int main(int argc, char *argv[]) {
    int n = atoi(argv[1]);
    for (int i = 0; i < n; i++) {
        int f = fork();
        if (f == 0) {
            printf("Child %d - PID %d - Parent %d\n", i, getpid(),
getppid());
            return 0;
        }
        else if (f < 0) {
            printf("Error creating child %d\n", i);
            return 1;
        }
    }
    return 0;
}
```

2.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

void f(int n) {
    if (n > 0) {
        if (fork() == 0) {
            printf("Child %d - PID %d - Parent %d\n", n, getpid(),
getppid());
            f(n - 1);
        }
        wait(0);
    }
    exit(0);
}

// create n child processes using fork and print pid of each child
int main(int argc, char *argv[]) {
    int n = atoi(argv[1]);
    //f(n);
    for (int i = 0; i < n; i++) {
        int f = fork();
        if (f == 0) {
            printf("Child %d - PID %d - Parent %d\n", i, getpid(),
getppid());
            wait(0);
        }
        else if (f == -1) {
            printf("Error\n");
        }
    }
}
```

```

        return 1;
    }
    else
        exit(0);
}
return 0;
}

```

LAB 13

```
#include <pthread.h>

pthread_mutex_ - init
    - lock
    - unlock
    - destroy
```

	R	W	
R	1	0	-> pthread_rwlock - init
W	0	0	- rdlock - wrlock - unlock - destroy

SEMAPHORES

```
#include <stdio.h>
#include <pthread.h>
#include <semaphore.h>
#include <unistd.h>

sem_t mutex;

void* thread(void* arg)
{
    //wait
    sem_wait(&mutex);
    printf("\nEnterd..\n");

    //critical section
    sleep(4);

    //signal
    printf("\nJust Exiting...\n");
    sem_post(&mutex);
}

int main()
{
    sem_init(&mutex, 0, 2);
    pthread_t t1,t2, t3;
```

```

pthread_create(&t1,NULL,thread,NULL);
sleep(1);
pthread_create(&t2,NULL,thread,NULL);

sleep(2);
pthread_create(&t3,NULL,thread,NULL);
pthread_join(t1,NULL);
pthread_join(t2,NULL);
pthread_join(t3,NULL);
sem_destroy(&mutex);
return 0;
}

```

SEMINAR 2

the return value of a function
 -> echo \$? to see it
 -> 0 if all went well

any shell program begins with #!/bin/bash
 -> reference to the version of shell i wanna use
 bash - born again shell :)

<- comment usually, but at the beginning it's a DIRECTIVE

if an argument doesn't exists, bash considers it an empty string
\$10 is not actually the 10th arg, rather the first one, to which we
append 10
for it to give me the 10th, \${10}

```
>./1.sh 1 2 3 4 5 6 7 8 9 0
First argument: 1; Second argument 2
The 10th argument: 0
```

\$* retruns all arguments as a string
\$@ retruns it as an array

when assigning variables, DONT USE WHITE SPACES, because it thinks
it's a command, and = is the arg

IFS=\$'\n' -> internal field separator

```
if condition; then
fi
```

anything that's not a link, pipe, directory is essentially a regular
file
for using grep to check if it's an int, we pipe it to it
> du -b -> tells you the numb of bytes of a file

```
1.sh
#!/bin/bash
```

```

# write a shell script that takes any amount of arguments and prints
if thaiti':as the name of a
# - regular file
# - directory
# - positive base 10 integer
# - anything else

echo "First argument: $1; Second argument $2"
echo "The 10th argument: ${10}"
echo "Argument count: $#" # argument count
echo "All arguments: $*" # returns it as a string
# echo "All arguments: $@" -> $@ returns it as an array
# echo "All arguments: $#" -> argument count
for a in $@; do
    var=0
    if test -f $a; then
        echo "$a is a regular file"
        var=1
    fi
    if test -d $a; then
        echo "$a is a directory"
        var=1
    fi
    if echo $a | grep -E -q "^[0-9]+$";then
        echo "$a is an integer"
        var=1
    fi
    if test $var -eq 0; then
        echo "$a is something else"
    fi
done

```

2.sh

```

#!/bin/bash
# write a shell script that takes as arguments pairs of arguments
# (filename, number)
# for each pair, check if the file size is less than the given num

size=0
while [ $# -eq 2 ]; do
    echo "args: $@"
    file=$1
    num=$2
    if [ -n $file ] && [ -f "$file" ]; then
        fsize=`du -s --block-size=1 $file | awk '[print $1]'`
        if [ $fsize -lt $num ]; then
            echo "File $file has size $fsize which is smaller than
$num"
        else
            echo "File $file has size $fsize which is greater than
$num"
        fi
        size=$((size + fsize))
    fi
done

```

```

    fi
    shift 2
done

echo "Total size: $size"

```

SEMINAR3

we've been working as if we're the only ones running on that machine
 usually a machine has 1 processor, 8 cores, 16 threads

ps - ef | wc -l -> see how many processes are running
 > ./1.sh in & -> runs the process in the background
 when two programs are running in the background, they are running in parallel, therefore you cannot predict the order of the output
 if those two processes are interacting with the same file or that interfere with one another, you might get a corrupted file
 > ./1.sh in & ./1.sh in & -> runs both the process in the background

moving forward, we'll assume that running two processes at the same time it's gonna mess our files

```
#!/bin/bash
```

```

COUNT=0
F=$1
while [ $COUNT -lt 100 ]; do
    X=$(cat $F)
    X=$((X+1))
    echo $X > $F
    COUNT=$((COUNT+1))
done

```

execl VS execlp

EXECL

for execl, you have to specify the full path of the file you want to execute

```
> which ls
/bin/ls
```

the initial argument is the name of the command you want to execute
 *remember int main(int argc, char ** argv)

*the list of arguments must be terminated by a NULL pointer
 execl("/bin/ls", "ls", NULL);

if you want to add more arguments, you just do it like this
 execl("/bin/ls", "ls", "-l", NULL);
 ^ first step into running another process in a C program

* exec replaces the current process with a new process, so it runs successfully, the current process will be terminated
 * in case exec fails, the current process will continue running
 -> this is called parasitic behaviour

```
! you should always check the return value of exec, see if it
failed or not
you can run whatever executable with execl
    execl("./1.sh", "1.sh", "./in", NULL);
```

EXECLP

```
execlp is looking for the path for us
    execlp("ls", "ls", "-l", NULL);
```

EXECV

```
-> not a variadic function, takes 2 arguments
execv(argv[1], argv + 1); -> ./exe /bin/ls -l
    -> this meets the criteria of the exec family having NULL as the
last argument
    * must have full path
```

EXECVP

```
execvp(argv[1], argv + 1); -> same, but full path is not needed
```

they all work in the same fashion, choosing which one to use is a matter of preference or situation

CTRL + C -> sends a signal to the process, SIGINT which iterrupts it
* man 7 signal -> see all signals available in linux
by calling the function signal() you can catch the signal and do something with it

```
signal(SIGINT, handler);
void handler(int sig) {
    printf("Caught signal %d\n", sig);
}
-> this will print "Caught signal 2" when you press CTRL + C
```

or signal(SIGINT, f) -> f is a function that will be called when the signal is caught

```
void f(int sig) {
    printf("Caught signal %d\n", sig);
}
signal(SIGINT, f);
```

because we've overwritten the default behaviour of the signal, we can't interrupt the program with CTRL + C anymore
now you have to SIGKILL the process

```
void f(int sig) {
    printf("Caught signal %d\n", sig);
    exit(0);
}
signal(SIGINT, f);
```

FORK()

```
-> creates a child process
-> everything after the fork() will be executed by both the parent and the child, essentially creating two processes
everything gets copied, the DS, CS, but also EIP (instruction pointer)
```

```

fork returns a value if it was successful
if (fork() < 0) {
    perror("fork");
    exit(1);
} else if (fork() == 0) {
    printf("Child PID = %d - PPID = %d\n", getpid(), getppid());
} else {
    printf("Parent PID = %d\n", getpid());
}
int i;
wait(&i);
-> this will make the parent wait for the child to finish executing

```

```

1.sh
#!/bin/bash

```

```

COUNT=0
F=$1
while [ $COUNT -lt 100 ]; do
    X=$(cat $F)
    X=$((X+1))
    echo $X > $F
    COUNT=$((COUNT+1))
done

```

SEMINAR5

```

up.c
/ create a thread for each argumentprintf("thrd %d - arg %s\n", d->id,
d->s);
// each thread converts any lowercase letters in the argument to
uppercase
#include <stdio.h>
#include <pthread.h>
#include <stdlib.h>
#include <string.h>
#include <ctype.h>

typedef struct {
    int id;
    char *s;
} data;

void* upcase(void *arg) {
    data *d = (data*)arg;
    char *copy = malloc(strlen(d->s) + 1);
    strcpy(copy, d->s);
    for (int i = 0; i < strlen(d->s); i++) {
        if (islower(d->s[i])) {
            copy[i] = toupper(copy[i]);
        }
    }
    printf("thrd %d - arg %s - copy %s\n", d->id, d->s, copy);
    return copy;
}

```

```

int main(int argc, char **argv) {
    int size = argc - 1;
    pthread_t th[size];
    data *args = malloc(size * sizeof(data));
    for (int i = 0; i < size; i++) {
        args[i].id = i;
        args[i].s = argv[i + 1];
        if (pthread_create(&th[i], NULL, upcase, &args[i])) {
            perror("Oh no");
        }
    }

    for (int i = 0; i < size; i++) {
        void *res;
        if (pthread_join(th[i], &res)) {
            perror("Error on join");
        }
        printf("Result from thread %d: %s\n", i, (char*)res);
        free(res);
    }
    free(args);
    printf("END: \n");
    for (int i = 0; i < size; i++) {
        printf("argv[%d] - %s\n", i+1, argv[i+1]);
    }
    return 0;
}

```

SEMINAR 6

`*((int *) v)` -> first casting, then dereference

we can have global variables

```

int n;
int *arr;

void *f(void *arg) {
    int index = *((int *) arg);
    arr[index] = rand() % 1000;

    // wait until all threads execute this line

    int other = rand() % n;
    while (other == index) {
        other = rand() % n;
    }

    int val = arr[other] / 10;
    arr[other] += val;
    arr[index] -= val;
    return NULL;
}

```

```
- BARRIER
    pthread_barrier_t bar;

    pthread_mutex_t mtx;

    void *f(void *arg) {
        int index = *((int *) arg);
        arr[index] = rand() % 1000;

        // wait until all threads execute this line
        pthread_barrier_wait(&bar);

        int other = rand() % n;
        while (other == index) {
            other = rand() % n;
        }

        // protect the shared variable when accessing it
        pthread_mutex_lock(&mtx);
        int val = arr[other] / 10;
        arr[other] += val;
        arr[index] -= val;
        pthread_mutex_unlock(&mtx);
        return NULL;
    }

int main() {
    srand(getpid());
    n = atoi(argv[1]);
    arr = malloc(n * sizeof(int));
    pthread_barrier_init(&bar, NULL, n);
    pthread_mutex_init(&mtx, NULL);
    pthread_t th[n];
    for (int i = 0; i < n; i++ ) {
        if (0 != pthread_create(&th[i], NULL, f, &i)) {
            perror("Cannot create thread");
            exit(1);
        }
    }

    for (int i = 0; i < n; i++) {
        pthread_join(th[i], NULL);
    }
    pthread_barrier_destroy(&bar);
    pthread_mutex_destroy(&mtx);
    for (int i = 0; i < n; i++) {
        printf("arr[%d] ", arr[i]);
    }
    free(arr);
    return 0;
}
```

```
!!! passing here i, allocated on the stack, always the same address,  
so we need to allocate it on the heap
```

```
    int index[n];  
    for (int i = 0; i < n; i++) {  
        index[i] = i;  
        if (0 != pthread_create(&th[i], NULL, f, &index[i])) {  
            perror("Cannot create thread");  
            exit(1);  
        }  
    }
```

```
! every time a barrier hits 0, it is reset to the initial value
```

- CONDITIONAL VARIABLES

```
pthread_cond_t cond;  
pthread_cond_wait (&cond, &mtx);  
    -> unlocks the mutex and waits for a signal  
    -> must be called with the mutex locked  
  
LOCK (mtx)           ----- Unlocks the mutex  
WHILE (condition) { | ---- thread waits for a signal  
    WAIT (cond, mtx); -----|  
                    | ---- thread is woken up  
                    | ----- unlocks the mutex  
}
```

-> if you call a signal with a locked mutex, you must unlock it before calling the signal

SEMAPHORE

```
sem_t sem1, sem2;  
sem_init(&sem1, 0, n/2); -> the flag in the middle is 0 for local, 1  
for global  
                           -> the last argument is the initial value of  
the semaphore  
sem_init(&sem2, 0, 1); -> is just a mutex  
  
sem_post(&sem1); -> increment the semaphore  
sem_wait(&sem1); -> decrement the semaphore
```

SEMINAR 7

Conditional variables

- one or more threads can wait for a condition to be true

pb 27 : Write a C program that takes two numbers, N and M, as arguments from the command line. The program creates N "generator" threads that generate random lowercase letters and append them to a string with 128 positions. The program will create an additional "printer" thread that waits until all the positions of the string are filled, at which point it prints the string and clears it. The N "generator" threads must generate a total of M such strings and the "printer" thread prints each one as soon as it gets to length 128.

```
shell exam
begin line: ^
end line: $
begin word: \<
end word: \>
```

2. write a regex that takes the multiple of 5 + 2 (ex: 2, 7, 12)
> grep -E "^(.{5})*..\$"
3. give the operators for the following shell verification:
 - a. string is not empty
[-n \$var]
 - b. argument is a regular file
[-f \$var]
 - c. is different (numerical)
[\$n1 -ne \$n2]
4. how would you restore the value of a file descriptor overwritten by dup2?
save the initial value of the file descriptor before calling dup2
5. 5
6. fork returns 0 or smt positive if there's a child

```
P
  /   \
C0   C1
  |   ^
  C01  C12 C13
 / \   |
C012 C013 C123
  |
C0123
=> 9 processes
```
7. only A

LECTURE 3
"in" file
"out" file
* daca o comanda se termina cu codul de return 0, atunci comanda este True (no error)

in shell, toate lucrurile se rezuma la codul ascii

```
constructia IF
$@ -> toate argumentele din linia de comanda, unul dupa altul
[ -f $A ] -> comanda test
#!/bin/bash
for A in $@; do
    if [ -f $A ]; then
        echo $A is a file
    elif [ -d $A ]; then
```

```

        echo $A is a dir
    elif echo $A | grep -Eq "^[0-9]+$" then
        echo $A is a number
    else
        echo We do not know what $A is
    fi
done

[ -z "$F" ] -> nu e vid
[ -r "$F" ] -> "citibil"

md5, sha -> o functie care primeste ca argument un fisier super mare,
care au ca output un sir de 160bits
hashing -> nu s-au gasit doua fisiere care sa aiba acelasi output
-> find > tipareste toate fisierele din acest director
-> [ -n "$STATE" ] lungimea nenula

*numele unui director, vreau sa anunt cand acesta a fost modificat
#!/bin/bash
D=$1 // numele directorului, primul argument
if [ -z "$D" ]; then
    echo "ERROR: No directory provided for monitoring" 1>&2 // iesirea
standard acelei comenzi echo, se va pune pe acealasi fisieri de erori
standard
    exit 1
fi
if [ ! -d "$D" ]; then
    echo "ERROR: Directory $D does not exist" >&2
    exit 1
fi

STATE=""
while true; do
    S=""
    CONTENT=""
    for P in `find $D`; do
        if [ -f $P ]; then
            LS=`ls -l $P | sha1sum`
            CONTENT=`sha1sum $P`
        else
            LS=`ls -ld $P | sha1sum`
            CONTENT=`ls -l $P | sha1sum`
        fi
        S="$S\n$LS $CONTENT"
    done
    if [ -n "$STATE" ] && [ "$S" != "$STATE" ]; then
        echo "Directory state changed"
    fi
    STATE=$S
    sleep 1
done

*anuntarea userilor cu prea multe procese
#!/bin/bash
```

```

maxim = $1
for user in `who | awk '{ print$1 }'`; do
    if [ 0`ps -u $user | wc -l` -ge 0$maxim ]; then # 0 daca cumva nu
se da $1
        echo $user
        write $user <<MESSAJ
        Aveti prea multe procese
MESSAJ
fi
done

*din directorul curent, sa se determine primul fisier text care
contine o linie al carei prim cuvant are cel mult 5 caractere
-> comanda file * - pentru directorul curent
-> putem decide ca un fisier este text daca comanda file
returneaza pentru el "ASCII text"
#!/bin/bash
    for x in * ; do
        if [ `file $x | grep -ci "ASCII text"` -eq 0 ] ; then continue
    ; fi
        cat $x | while read cuv1 t; do
            if [ ! -z $cuv1 ] && [ `expr length $cuv1` -ge 5 ]; then
                echo In $x s-a gasit $cuv1 cu lungimea `expr length
$cuv1`
                break
            fi
        done
        break
    done

```

LECTURE 4

* is anyone gonna mess with my data while i use it -> the new 2AM formula

ex:

```

load AX, n
inc AX
store AX, n

```

```

A | B
L
    L
I
    I
S
    S

```

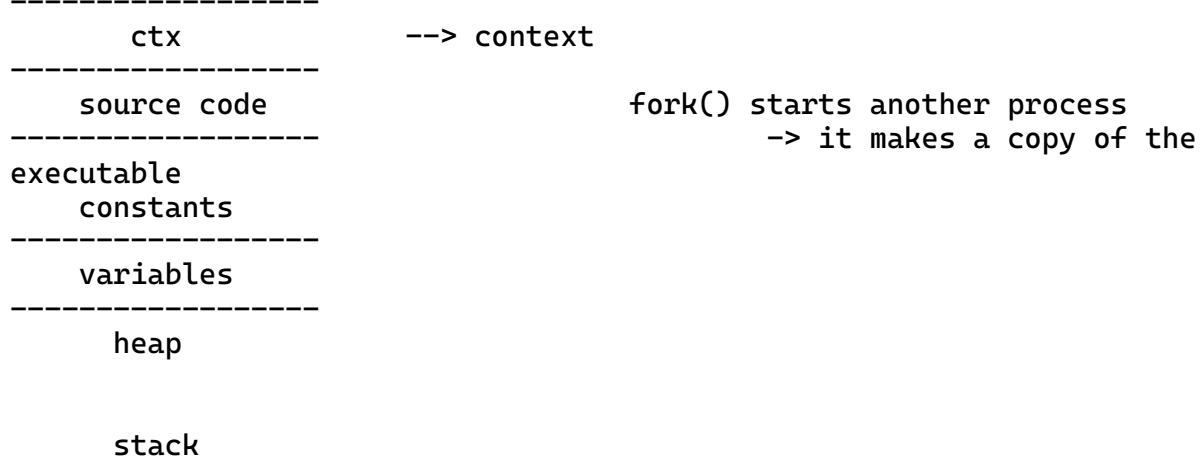
this way the data is compromised, since it happened at the same time. your value was not incremented

& at the end of a line runs it in the background

how do you have multiple programs running on the same processor at the same time?

-> a program is a file in execution
-> every few seconds, the processor interrupts the execution and chooses another program, puts it onto the CPU

bash knows how to create a process, all commands are essentially that :)



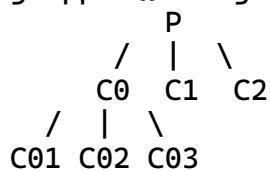
-> output of 1.c
#include <stdio.h>
#include <unistd.h>

int main(int argc, char** argv) {
 printf("before %d %d\n", getpid(), getppid());
 fork();
 printf("after %d %d\n", getpid(), getppid());
 return 0;
}
before 12173
after 12173
after 12174

before 12199 11884
after 12199 11884
after 12200 12199

getpid() -> get the id of the child program

getppid() -> get the id of the parent program



-> for a for loop i=0,3

```

before 12251 11884
after 12251 11884
before 12251 11884
after 12251 11884
before 12251 11884
after 12252 12251
before 12252 12251
after 12253 12251
before 12253 12251
after 12251 11884
after 12252 12251
before 12252 12251
after 12254 1
after 12255 12252
before 12255 12252
after 12253 1
after 12252 1
after 12256 1
after 12257 12252
after 12255 1
after 12258 1

* fork() returns
-> 0 in the child
-> child PID in the parent

-> output for 2.c
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    int pid;
    pid = fork();
    printf("before\n");
    if (pid == 0) {
        printf("Child specific code pid=%d for fork return=%d.\n",
getpid(), pid);
        exit(0);
    }
    printf("parent specific code child PID=%d.\n", pid);
    wait(0);
    return 0;
}

before
parent specific code child PID=12546.
before
Child specific code pid =12546 for fork return=0

wait() does not specify which child to wait for
-> it's job is to wait for the children processes to finish
* all children need to die before the parent finishes the job

```

```

-> if google were to be implemented like this, only one process could
be done at a time
while(1) {
    get request
    process request
    send back response
}

-> in this version, getting the request is the only thing that's done
one by one
while(1) {
    get request
    if (fork() == 0) {
        process request
        send response
        exit(0);
    }
    // wait(0);
}

```

zobies -> children that don't die, because wait is not called
you cannot really add a wait() here but it's gonna be very
inefficient. since it gets a request, but does nothing with it until
all children processes finish

LECTURE 5

children agian??

```

while(1) {
    get req
    if (fork() == 0) {
        process req
        send response
        exit(0);
    }
}

```

when hitting ctrl c to interrupt the prgram, that's a signal
same goes for ctrl z which places it in the background

with few exceptions, all signals stop the program
* man signal

make a program that refuses to stop
we'll use 2 function, signal and kill
kill sends signals, signal stores it
TO KILL THE PROCESS
> kill -SIGINT 8436
> kill -9 8436
kill is a function that takes the signal and the runtime (?) id

both wait and sleep stop the cpu
wait on the other hand waits for the child process to finish

```

void f(int signal){
    wait(0);
}

signal (SIGCHLD, f);
while(1) {
    get req
    if (fork() == 0) {
        process req
        send response
        exit(0);
    }
}
SIG_IGN -> ignores the signal
SIG_DEF -> default

```

Exec*

	absolute path	name to be searched in path
array	execv	execup
list	execl	execlp

array	execv	execup
list	execl	execlp

```

export PATH=.:$PATH
^ this doesn't change the path, it prefixes it
export PATH=.
^ this changes the path

execlp("grep", "grep", "-E", "abc", "/etc/passwd")
char *a[] = {"grep", "grep", "-E", "abc", "/etc/passwd", NULL}
execup("grep", a);

*rewrites echo
#include <stdio.h>
#include <unistd.h>

int main(int argc, char** argv) {
    execlp("echo", "echo", "asdf", NULL);
    printf("we echoed\n");
    return 0;
}

creating links between files: > ln -s ./a x
specify the program you want to run twice when using execlp!!

```

1	2	3	4	5	6
-----	-----	-----			
3		7		11	
-----		-----			

10

21

^ how to add processes fast
-pipe
-fifo
-ipc -> shared memory (SHM)
 -> message queue (MSG)
 -> ?? smt (SEM)

PIPE

> ALWAYS CLOSE THE PIPE ENDS AS SOON AS YOU DON'T NEED THEM
> close() doesn't actually close the pipe, but rather the pipe
handlers

LECTURE 7

read

- > synchronisation process, not your usual read/write
- extracts the data (doesn't leave it in place)
- waits when empty (doesn't give an error either)
 - until some data
 - until no writers
- * read doesn't give you what you want, but rather what it can

write

- adds data (might overwrite)
- waits when full
 - until some space
 - until no readers

* none of them gives an error
-> Why doesn't read read as many bytes as i ask it to?
 maybe there aren't that many
 how long are you gonna wait for it?

open (FIFO)

- waits until FIFO opened for complementary operations
- O_NDELAY (fcctl) -> flag that can be passed to open and it does NOT wait for anything

pipes are ONE DIRECTION!!

pipes and FIFO are communication channels

IPC - inter processor communication (SHM, MSG, SEM)

-> mechanisms at the processor level that allow communication between them

when you create this, try the ip to be unique

-> ipcmk

- clean up after yourself :)

> ipcmk --shmem 1024 --mode 600 (make a shared memory)

> ipcs (see it)

```

> ipcrm -M [key] (remove it)

how does it work?
-> create it
-> where you open it, you get a handler

shmget has no synchronization, it doesn't wait!
by the time b is ran, you have no more shared memory either

h.h
#include <stdio.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#include <stdlib.h>
#include <time.h>

struct absp {
    int a;
    int b;
    int s;
    int p;
};

a.c
#include "h.h"

int main(int argc, char ** argv) {
    int shmid;
    struct absp* x;

    shmid = shmget(1234, sizeof(struct absp), IPC_CREAT|0600); // 
create a shared memory at ip 1234 and give it permission 0600
    x = shmat(shmid, 0, 0); // pointer to that memory, shmat just
connects it
    srand(time(NULL)); // current time in seconds since 1970 :)
    while (1) {
        x->a = rand() % 100;
        x->b = rand() % 100;
        if (x->p == x->s)
            break;
    }

    shmdt(x);
    shmctl(shmid, IPC_RMID, NULL);
    return 0;
}

b.c
#include "h.h"

int main(int argc, char ** argv) {
    int shmid;

```

```

    struct absp* x;

    shmid = shmget(1234, 0, 0);
    x = shmat(shmid, 0, 0); // pointer to that memory
    srand(time(NULL)); // current time in seconds since 1970 :)
    while (1) {
        x->p = x->a * x->b;
        x->s = x->a + x->b;
        if (x->p == x->s)
            break;
    }

    shmdt(x);
    return 0;
}

```

threads are mechanism that take a function and runs them WITH the current process, in parallel

- > unlike fork(), you don't clone the whole process
- when creating a thread, you create a new stack

 ctx

 source

 const

 variables

 heap

 stack

THREAD

```

#include <stdio.h>
#include <pthread.h>

void* f(void* a) {
    printf("f\n");
    return NULL;
}

int main(int argc, char ** argv) {
    // threads will have a handler, but there's no child-parent
    relation here
    pthread_t t;
    pthread_create(&t, NULL, f, NULL);
    // now we have 2 threads, main and f that work in parallel
    printf("main\n");
    pthread_join(t, NULL);
    return 0;
}

```

one argument is enough since you can pass a struct with all the arguments you need

LECTURE 8
-race condition
 $n \rightarrow$ critical
 variable
 resource
 $n++ \rightarrow$ critical region

mutex (mutual exclusion)
rwlock (read-write lock)
contional variable
barriers
semaphores

! surround critical region with mutex lock and unlock

LECTURE 10

prb 1: you have a bunch of kids trying to go into a bathroom stall
 $\text{trylock}() \rightarrow$ either locks or returns a value

\rightarrow problem with kids and bathroom stalls
 for (int i = 0; i < 3; i++) {
 lock(&m[i]);
 1 go
 unlock(&m[i]);
 }

\rightarrow this will make everyone wait for the first stall, not
distribute them

prb 2: elevator \rightarrow more kilos than allowed, the elevator keeps locking

prb 3: trampoline \rightarrow only a specific number of kids can jump at a time

prb 4: forest \rightarrow bees and bears
 \rightarrow bees make honey, bears eat honey
 \rightarrow ranger kicks in when there's no honey

forest.c
#include <stdio.h>
#include <pthread.h>
#include <seamaphore.h>
#define BEES 1
#define BEARS 5

```
int honey = 100;
pthread_mutex_t m;
pthread_cond_t c;
sem_t s;

void* bear(void* a) {
    while (1) {
        pthread_mutex_lock(&m);
        if (honey >= 10) {
            printf("-");
            honey -= 10;
        }
    }
}
```

```

        }
    else {
        printf("!");
        pthread_cond_signal(&c);
    }
    pthread_mutex_unlock(&m);
}
return NULL;
}

void* bee(void* a) {
    while (1) {
        pthread_mutex_lock(&m);
        printf("+");
        honey++;
        pthread_mutex_unlock(&m);
    }
    return NULL;
}

void* ranger(void* a) {
    while (1) {
        pthread_mutex_lock(&m);
        while (honey >= 10) { doesn't go straight to producing honey, but rechecks
            pthread_cond_wait(&c, &m); // wait -> unlock | waits for
signal | lock
        }
        printf("H");
        honey += 100;
        pthread_mutex_unlock(&m);
    }
    return NULL;
}

int main(int argc, char** argv) {
    int i;
    pthread_t bees[BEES], bears[BEARs], rangert;
    pthread_mutex_init(&m, NULL);
    pthread_cond_init(&c, NULL);
    sem_init(&s, 0, 0);
    for (i = 0; i < BEES; i++) {
        pthread_create(&bees[i], NULL, bee, NULL);
    }
    for (i = 0; i < BEARs; i++) {
        pthread_create(&bears[i], NULL, bear, NULL);
    }
    pthread_create(&rangert, NULL, ranger, NULL);
    for (i = 0; i < BEES; i++) {
        pthread_join(bees[i], NULL);
    }
    for (i = 0; i < BEARs; i++) {
        pthread_join(bears[i], NULL);
    }
    pthread_join(rangert, NULL);
}

```

```
pthread_mutex_destroy(&m);
pthread_cond_destroy(&c);
sem_destroy(&s);
return 0;
}
```

LECTURE 11

Operating Systems Structure



KERNEL

Tech book keeping

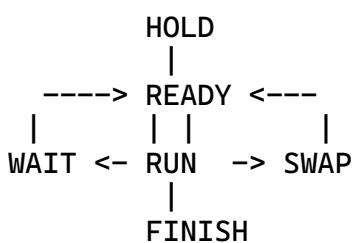
```
    interrupts
    / (talk to the hardware)
hardware - process manager      -> memory management -> phisical
I/O -> File System -> Job planning & resource allocation
          \ (chosses processes)
            CPU dispather
              (?)
```

Book keeping

```
Compiler + Assemblying + Linkers + Loaders
Interpreters + Macro processors + Editor + Debuggers      -> IDEs
Utils + Libraries + DataBases + Other
```

Processes

states



Scheduling

- FCFS (FIFO order)
- SJF (Small Jobs First)
- Priorities
- Deadline scheduling

DeadLock

LECTURE 12

Conditions that make dead lock possible

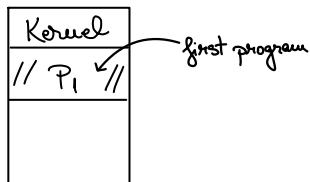
1. mutual exclusion
2. lock & wait
3. non-preemption
4. circular wait

ALWAYS LOCK THE RESOURCES IN THE SAME ORDER
+ open files in the same order
+ make sure you don't have cycles
+ sum capacity is less than all threads (sumo elevator problem)
+ barrier is never completed \Rightarrow no. of threads is not a multiple of the barrier init

Memory management

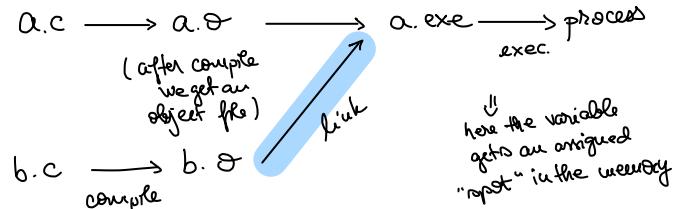
- real \rightarrow single tasking OS (A)
 \rightarrow multitasking OS
 - fixed partition * absolute (B)
 - * relocatable (C)
- variable partitions (D)
- virtual
 - \rightarrow paged (E)
 - \rightarrow segmented (F)
 - \rightarrow paged - segmented (G)

A. Single tasking OS



-the compiler hard codes physical memory addresses in the executable

Address calculation:

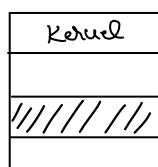


+ there's a tight connection between the .exe and the machine, since the compiler knows exactly how big the kernel is

!! here the variable gets an original "spot" in the memory

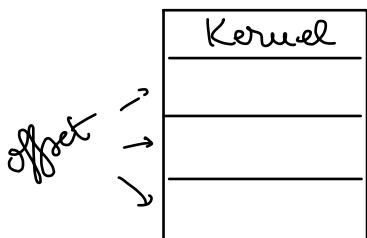
* stopped working for us once we were running more programs

B.



- \rightarrow split into memory partitions
- \rightarrow compile into certain partition

C. Change compiler and improve address computation



- no hardcoding of physical addresses
 - every partition has an offset
 - relative addresses (relative to the beginning of the partition)
 - the program can be in any partition (that is free)
- * $\text{Dinony address} - \text{offset}$
 $\text{Partition address} = \text{partition start} + \text{DA}$

D partitions are not predefined, we define them based on the size of the program that has to start (each new process starts after the last one)

Kernel
// P ₁ //
// P ₂ //
// P ₃ //
// P ₄ //
// P ₅ //



Kernel
// P ₁ //
// P ₂ //
// P ₃ //
// P ₄ //
// P ₅ //

we can run P₆ only if we have the continuous memory necessary



We have the necessary memory but it is SPLIT ⇒ fragmentation

RAM

		5	6
2			
		1	
	7	8	
4			
0			3

↑
real pages (fixed size of a page)

Prog

0	1	2
3	4	5
6	7	8

↑
virtual pages

- when the process stops, the real pages are freed
- fragmentation solved with more fragmentation
- more complex address calculation
needs a search

→ we need a table for each virtual page position in the real pages ⇒ slow

- virtual address = virtual page + offset
- ⇒ we fixed fragmentation

BUT the processor help with hardware

F. segments do not solve fragmentation, segments only group sections to be protected

they do not have a fixed size | provides memory access protection
not meant to organize your code + address calculation similar to page

G. TOP but waste a bit more memory than just paged

every segment split into pages

→ address : segment table + page table → virtual address (s, rp, offset)

Loading policies:

→ when should pages be loaded?

- load all of them at process start
(slower start and wasted RAM) loading pages that might never be used
- load when needed (a location from it is referred)
(even slower start? and slower execution)
- load the requested page and a few neighbouring pages
(chances are they will be needed)
Neighbouring principle (if it requests a page, it is likely to soon request its neighbouring pages)
↳ what we use today

Unloading policies:

RAM

	0		
3			1
	2		

Prog:

0	1
2	3

LECTURE 13

Unloading policies / replacement policies

-> RAM memory

RAM	Program
- - - -	0 1
- 0 - -	2 3
3 - - 1	
- - - -	
- 2 - -	

* choose something unlikely to be needed soon

-> What if the memory is full?

- you spill some data to disk

-> How to decide what to spill?

? FIFO - little to no prediction (what if the first one is a c library?) or the printing to screen

- NRU (Not Recently Used) - use 2 bits to track the usage of each page (that are periodically set to 0)

1 - 00 - not used recently, not modified, not read, not written

3 - 01 - not used recently, modified, not read, written

2 - 10 - used recently, not modified, read, not written

4 - 11 - used recently, modified, read, written

-> you use the file from where you read more often than the one you write to * Choose a victim from the smallest class

LRU - last recently used

N pages of RAM

NxN bits matrix

when a page is accessed, its line is filled with 1, it's column with 0

the victim is the page with the minimum line total

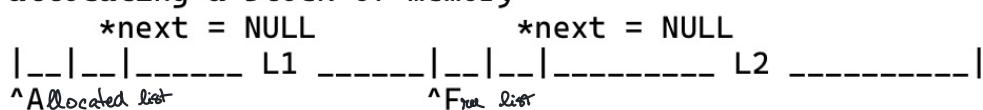
0 0 0 0	0 0 0 0	0 0 0 0	0 0 0 0
0 0 0 0	1 0 1 1	1 0 1 0	1 0 0 0
0 0 0 0	1 0 0 0	0 0 0 0	1 1 0 1
0 0 0 0	0 0 0 0	1 1 1 0	1 1 0 0

Allocation policies / Placement policies

----- where do we malloc?



allocating a block of memory



- Where do you place a malloc request?

-> First Fit (choose the first place big enough)

- fast

- but you make no attempt to minimize fragmentation

-> Best Fit (choose the tightest fitting space)

- slower

- very small slices of free memory are created and left behind

-> Worst Fit (allocate the largest one available)

- slow

fine grained fragmentation

- Buddy Fit
- leaves the largest possible chunks of free memory
- $$2^n = 2^0 + 2^0 + 2^1 + 2^2 + \dots + 2^{(n-1)}$$
- ↳ always allocate a multiple of 2 / keep lists for each power of 2 of free spaces / allocate the closest / split the remaining list across the other list
- What's the point of the cache?
- to reduce the time it takes to access data
 - you keep stuff that you use often in the cache, close to you

Memory Hierarchy

Registers

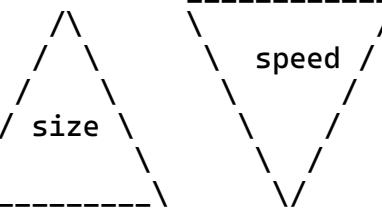
L1 cache

L2 cache

L3 cache

RAM

Disk (HDD/SSD)



Cache → small + fast
size \leq pages

RAM → large + slow
size \geq pages

Direct swapping / Direct cache

cache location = page number % cache size

cache size = 16

0 8 17 1 33 49 2

0 8 1 1 1 1 2

* you cancel the cache's usefulness, since there's collision

Associative swapping

place the page in the first empty slot

* organize cache in groups of pages, choose a group with %, search for page / free spot within the group

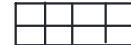
Set Associative swapping

0 | 0 8 2 . . .

$C=8$

0 1 2 3

$k \rightarrow k \% 4$ and search



File Systems

- data is stored in blocks
- file - a sequence of blocks

I-Node

| inode data | 0 | 1 | . . . | 9 | 10 | 11 | 12 |

*pdf mara

Directory

Name	I-Node
------	--------

a.txt	17
b.c	23
dir/	108

* poza links? ask mara

Remounting Volumes

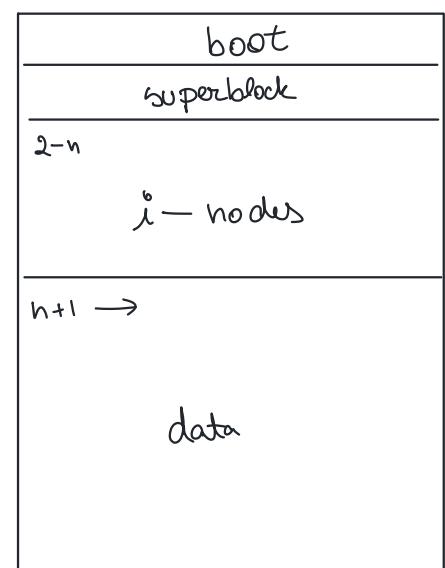
/: | A | B |

\

B: | X | Y |

stack

/: | X | Y |



* how large can a file be?

if: address size A

block size $N * A$

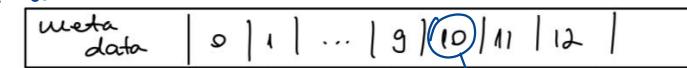
$$\Rightarrow 10 * N * A + N^3 * A - N^2 * A$$

* each file has an i-node as its entry

* a directory is just a file that contains pointers to other files

UHIX FS: → data is stored in blocks // interfațare între sistemul de fișiere și
 (file system) file = collection of blocks locația efectivă

J-mode → keeps track of where the data is
 What an inode looks like ← addresses →



↑
 file size
 permissions
 ownership

↓
 0
 0

↓
 0

↓
 0

points to a block
 that points to data

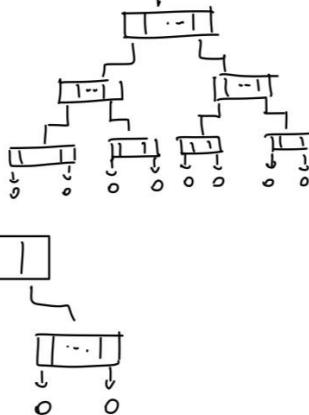
they go directly
 to a new block

single
 redirection

double
 redirection

A = address size
 B = block size

triple
 redirection



- * se poate face legătură hard dacă face parte din același partit?
- * diferența dintre un link simbolic și un link hard?
- o legătură simbolică nu e o copie a fișierului
 - linkul e ca un "pointer" spre fișier
 - dacă stergem fișierul original, un link soft nu mai are speră ce să point
 - DAR un link hard e okay, pentru că e copie a mode-ului
 - un link soft pointerii spre mode-ul fișierului
 - ln -s [fișier] → soft link
- * dacă modificăm link-ul hard se modifică și fișierul de pe disk
- * ce conține un fișier de tip director în Linux?

Fișierele director. Un fișier **director** se deosebește de un fișier obișnuit numai prin informația conținută în el. Un **director** conține lista de nume și adrese pentru fișierele subordonate lui. Uzual, fiecare utilizator are un **director** propriu care punctează la fișierele lui obișnuite, sau la alți subdirectorii definiți de el.

Fișierele speciale. În această categorie putem include, pentru moment, ultimele 6 tipuri de fișiere. În particular, Unix privește fiecare dispozitiv de I/O ca și un fișier de tip special. Din punct de vedere al utilizatorului, nu există nici o deosebire între lucrul cu un fișier disc normal și lucrul cu un fișier special.

Fiecare **director** are două intrări cu nume speciale și anume:

" . " (punct) denumește generic (punctează spre) însuși **directorul** respectiv;

" .. " (două puncte successive), denumește generic (punctează spre) **directorul** părinte.

Fiecare sistem de fișiere conține un **director** principal numit **root** sau **/**.

* deadlock

↳ in process some thread accepts a resource so file
elaborate de un proas, care la ramand lui accepte

```
fifo1 = open("f1", O_RDONLY);  
fifo2 = open("f2", O_WRONLY);  
[...]
```

```
fifo2 = open("f2", O_RDONLY);  
fifo1 = open("f1", O_WRONLY);  
[...]
```

process 1 accepts dup 2 si invins

sem

last writer si locked la scriere din thread

* signal

↳ handler + cum se înregistrează

↳ semaphore care nu pot fi regizate

Signal (SIGINT, handler);

↳ asociare interrupt-ul cu functia handler.

Producer / consumer problem

threads communicating through pipe

| want the producer to wait if buffer is full
| & consumer to wait if buffer is empty

2 semaphores: full(0) empty(N)

P	C
wait(empty) produce post(full)	wait(full) consume post(empty)
	lock(w) unlock(w)

to see the difference between 2 directories you can use
diff goes line by line and cmp goes byte by byte

```
find "$1" -type f | while read F; do
```

```
    find "$1" -type f | while read G; do
```

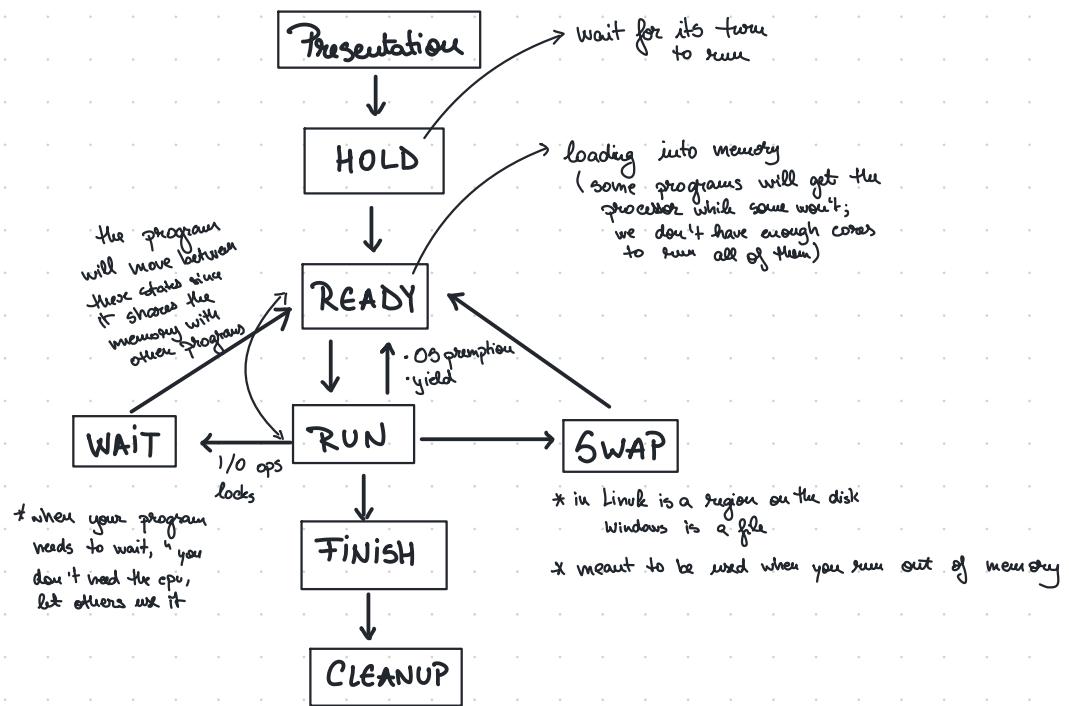
```
        if [ "$F" != "$G" ] && cmp -s "$F" "$G";  
            echo "$F = $G"
```

```
done done fi
```

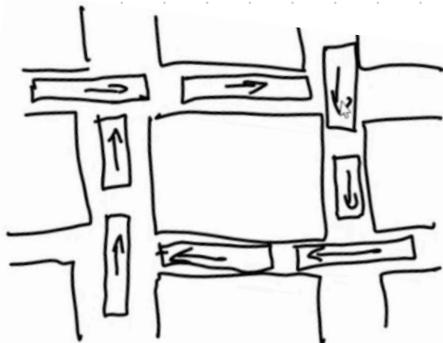
Lecture 12

Process Management

→ process states



Deadlock



Only solution is to choose a victim to back up

like:

A

lock(x)

lock(y)

B

lock(y)

lock(x)

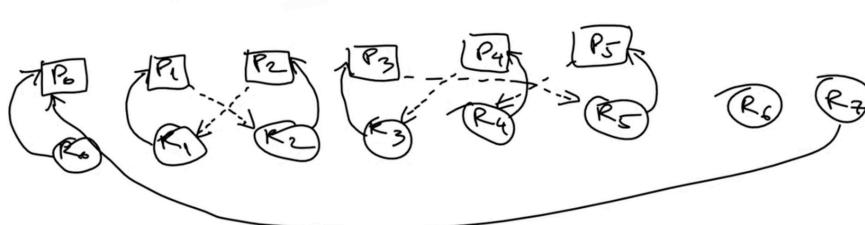
the invisibility creates the deadlock

* getting out of a deadlock?

- stop one of the processes / threads

- maybe (if you have the tech) rollback to a saved safe state

* detect deadlock



Legend

- [] - process
- - resource
- - locked
- > - waiting on a lock

deadlock 1

deadlock 2

? Cycles

if you detect a cycle, that's a deadlock

* preventing a deadlock

what makes it possible?

1) Mutual exclusion (resources are how to stay)

2) Hold (lock) and wait (not reusable)

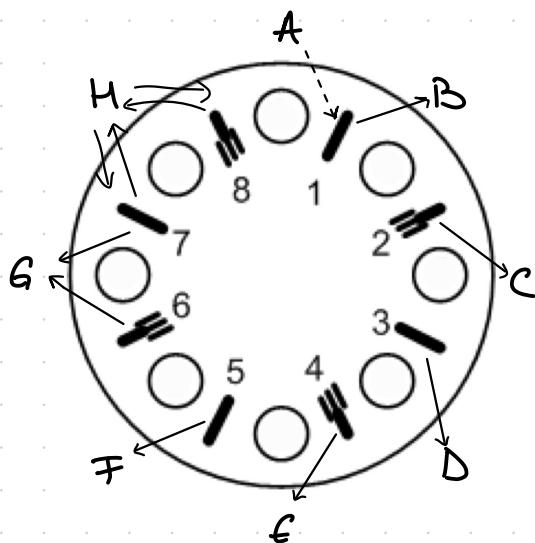
3) Non-preemption (someone else won't be able to steal the lock)

4) Circular wait

(break by locking in the SAME order)

This is how you prevent a deadlock

Philosopher's problem



to eat you need both sticks

if everyone picks the one on the right we all eats
⇒ deadlock

? impose to choose the same order of sticks all the time
ex: pick up the stick with the lower number

Process scheduling

- FCFS - first come first serve

- SJF - shortest job first

→ client will have to provide an estimate of the execution duration

* causes starvation, the guy will more time will never get to the CPU

- use priorities * mars robot story when priorities fail

- deadline scheduling

- Round Robin - assign small quanta of time to each process

Memory management

1) Allocation policies

2) Replacement policies

3) Placement policies

- Real allocation

- single tasking systems
- multi tasking systems
 - fixed partitions
 - absolute
 - relocatable
 - dynamic partitions