

1. ANSI-SPARC architecture:
  - a. Conceptual structure (schema) - just 1
  - b. External structures (ex. views) - there can be several
  - c. Physical (internal) structure (data files, indexes etc.) - just 1 *⇒ why we can have just one clustered index*
  - d. No such thing as a symbolic structure
2. **Data independence** -
  - a. Logical - apps + external schemas are not affected by changes in the conceptual structure
  - b. Physical - apps + ext sch are not... physical structure
3. From the simplest to the most complex: field, record, table, database
4. Data description model can be used to describe: the structures of the data, relationships with other data, consistency constraints
5. The rows in a relational table are not ordered, the records are distinct - but DBMSs allow duplicates
6. Degree of a relationship set = the no. of entity sets the participate in that rel
  - a. The degree of a projection < .. of the original table
7. **Optimizer** - produces an efficient execution plan for query eval, based on storage information
8. Multiple candidate keys can be declared using UNIQUE; one of them can be chosen as the primary key
9. We can have a multiple FK that references a multiple PK
10. FK can be null
11. An att can be a PK and FK at the same time
12. **Key** = a restriction defined on an entity set; a set of attributes with distinct values in the entity's set instances
13. **Integrity constraints** = conditions specified on the DB schema, restricting the data that can be stored in the DB (includes key constraints, FK constraints)
  - a. Key constraint = a constraint stating that a minimal subset of atts in a relation is an unique identifier for every tuple in the relation
  - b. **Superkey** = set of fields that contains the key
14. **3-valued logic** (true, false, unknown = null)
15. "expression = ANY (subquery)" ⇔ "expression IN (subquery)"
16. "Expression <> ALL (subquery)" ⇔ "expression NOT IN (subquery)" *different*
17. Copy data from one table to another: "INSERT INTO T2; SELECT \* FROM T1"
18. **SELECT** can contain arithmetic operations
19. **SELECT**, **FROM** - mandatory; **HAVING** - optional
20. **HAVING** cannot contain row-level conditions because it works with groups

21. Patterns: "%" - 0 or more arbitrary characters, "\_" - any one character (used with LIKE)
22. Can't group by a subquery
23. SELECT R.\*  $\Leftrightarrow$  select all the columns in R, ex.  
SELECT R.\*  
FROM R  
RIGHT JOIN whatever ON whatever
24. LEFT JOIN = inner join + the left operands which don't match the condition (their columns are filled with NULL); same idea for RIGHT JOIN
25. If we have a FROM query, we need to give it an alias
26. Processing order: from, where, group by, having, select, distinct, order by, top
27. **Functional dependency**: left side = **determinant**; right side = **dependent**
28. Functional dep properties:
  - a. K - key of R,  $\mathcal{P}$  - key of a subset of R  $\Rightarrow K \rightarrow \mathcal{P}$
  - b.  $\mathcal{P} \subseteq \alpha \Rightarrow \alpha \rightarrow \mathcal{P}$  (trivial func. dep., reflexivity)
  - c.  $\alpha \rightarrow \mathcal{P} \Rightarrow \gamma \rightarrow \mathcal{P}$ , for any  $\gamma$   $\alpha \subset \gamma$
  - d.  $a \rightarrow b$  and  $b \rightarrow c \Rightarrow a \rightarrow c$  (transitivity)
  - e.  $a \rightarrow b \Rightarrow a \cup c \rightarrow b \cup c$  (augmentation)
29. **Prime attribute** if it's included in a key K; non-prime otherwise
30. B is **fully functionally dependent** on A if  $A \rightarrow B$  and **B doesn't depend on any proper subset of A** (B cannot be derived from a smaller part of A)
31. **Multivalued dependencies** - each value from the determinant is associated with a set of values, not with just one, denoted by " $\rightarrow$ " " $\Rightarrow$ "
32. **Join dependency** (JD) - we project an R[A] onto  $R_1[\alpha_1] \dots R_n[\alpha_n] \Rightarrow R$  satisfies the JD  $\{\alpha_1, \dots, \alpha_n\}$  if  $R = R_1 * \dots * R_n$
33. **Normal forms**:
  - a.  $1NF \subset 2NF \subset 3NF \subset BCNF \subset 4NF \subset 5NF$ 
    - i. If not 1NF  $\Rightarrow$  not 2,3..NF; if not 2NF  $\Rightarrow$  not 3,BC..NF and so on
  - b. 1NF = no repeating attributes
  - c. 2NF = 1NF + **every** non-prime att is fully functionally dependent on every key of the relation
  - d. 3NF if for any non-trivial  $X \rightarrow A$ , X is a superkey or A is a prime att
  - e. (Boyce-Codd) **BCNF** if every determinant (left side of a func dep) is a key
    - i. Most desirable form
  - f. 4NF if for every multi-valued dep  $a \twoheadrightarrow b$ , either one is true
    - i.  $b \subseteq a$  or  $a \cup b = R$   
*cover all*

- ii.  $a$  - superkey
- g. 5NF = every non-trivial join dependency (JD) is implied by the candidate keys in  $R$ 
  - i. A JD  $\{ \alpha_1, \dots, \alpha_n \}$  is trivial if any  $\alpha_i$  is the set of all attrs of  $R$
  - ii. ... is implied by the candidate keys in  $R$  if ALL  $\alpha_i$  are superkeys in  $R$

34. **Armstrong's axioms:** reflexivity, augmentation, transitivity

35. Rules derived from AA

- a.  $a \rightarrow b$  and  $a \rightarrow c \Rightarrow a \rightarrow b \cup c$  (union)
- b.  $a \rightarrow b \cup c \Rightarrow a \rightarrow b$  and  $a \rightarrow c$  (decomposition)
- c.  $a \rightarrow b$  and  $b \cup c \rightarrow d \Rightarrow a \cup c \rightarrow d$  (pseudotransitivity)

36.  $\alpha^+ =$  closure of  $\alpha$  under a set of fds

37. **Example: Show that some fds are in  $\alpha^+$**

We have a relation  $R = \{a \dots f\}$  and  $\alpha = \{a \rightarrow b, a \rightarrow c, cd \rightarrow e, cd \rightarrow f, d \rightarrow e\}$ . Show that the following fds are in  $\alpha^+$

$a \rightarrow bc$ : union of  $a \rightarrow b$  and  $a \rightarrow c$

$cd \rightarrow ef$ : union of  $cd \rightarrow e$  and  $cd \rightarrow f$

$ad \rightarrow e$ :  $a \rightarrow c$  / augment with  $d \Rightarrow ad \rightarrow cd$ ; from  $cd \rightarrow e$  through transitivity  $\Rightarrow ad \rightarrow e$

$ad \rightarrow f$ : pseudotransitivity  $a \rightarrow c$  and  $cd \rightarrow f$  or same as before, augment first

38. Algorithm for computing  $\alpha^+$

**closure** :=  $\alpha$ ;

**repeat until there is no change:**

**for every functional dependency  $\beta \rightarrow \gamma$  in  $F$**

**if  $\beta \subseteq$  closure**

**then closure := closure  $\cup \gamma$ ;**

39. **Example: Compute  $\alpha^+$**

$R = \{a \dots f\}$  and  $F = \{a \rightarrow b, a \rightarrow c, cd \rightarrow e, cd \rightarrow f, d \rightarrow e\}$

$\alpha = \{a, d\}$  (given)

$\alpha^+ = \{a, d\}$

$a \rightarrow b \Rightarrow \alpha^+ = \{a, b, d\}$

$a \rightarrow c \Rightarrow \alpha^+ = \{a, b, c, d\}$

$cd \rightarrow e \Rightarrow \alpha^+ = \{a, b, c, d, e\}$  (because both  $c, d$  were in  $\alpha^+$ )

$cd \rightarrow f \Rightarrow \alpha^+ = \{a, b, c, d, e, f\}$  (same reason)

$d \rightarrow e \Rightarrow$  we don't add it because  $E$  already is in  $\alpha \Rightarrow$  no change at this step

$\Rightarrow$  stop  $\Rightarrow \alpha^+ = R$

40.  $F, G$  - sets of fds  $\Rightarrow F \equiv G$  (equivalent) if  $F^+ = G^+$
41. **Minimal cover**  $F_m$  if
- $F_m \equiv F$
  - Right side of every dependency in  $F_m$  has a single att
  - Left side .. is irreducible (if we remove an att from it, we modify  $F_m$ 's closure)
  - No dependency in  $F_m$  is redundant (if we remove it, we modify the closure)
42. **Example: compute  $F_m$**
- $R = \{a \dots d\}; F = \{a \rightarrow bc, b \rightarrow c, a \rightarrow b, ab \rightarrow c, ac \rightarrow d\}$
- Use decomposition on  $a \rightarrow bc \Rightarrow a \rightarrow c, a \rightarrow b$  (which already was in, so we don't add it again)  $\Rightarrow \{a \rightarrow b, a \rightarrow c, b \rightarrow c, ab \rightarrow c, ac \rightarrow d\}$
- $a \rightarrow c$  / augment with  $a \Rightarrow aa \rightarrow ac \Rightarrow a \rightarrow ac$ ; we know  $ac \rightarrow d \Rightarrow a \rightarrow d$  (transitivity); now  $ac \rightarrow d$  becomes redundant, we replace it with  $a \rightarrow d \Rightarrow \{a \rightarrow b, a \rightarrow c, b \rightarrow c, ab \rightarrow c, a \rightarrow d\}$
- $a \rightarrow c$  / augment with  $b \Rightarrow ab \rightarrow cb \Rightarrow$  (decomposition)  $ab \rightarrow c$  (it's already in, we can remove it)  $\Rightarrow \{a \rightarrow b, a \rightarrow c, b \rightarrow c, a \rightarrow d\}$
- $b \rightarrow c$  can be obtained as  $a \rightarrow b$  and  $b \rightarrow c$  (transitivity)  $\Rightarrow$  remove  $b \rightarrow c$
- $F_m = \{a \rightarrow b, a \rightarrow c, a \rightarrow d\}$
43. Just by looking at an instance, we can **only** say that a func. dep is **satisfied** or not, we can't conclude that it's **specified on the schema**
44. Algebra symbols
- $\sigma_C(R)$  = select \* from R where C; distributive over set diff, intersection, union
  - $\pi_\alpha(R)$  = projection; select (distinct?)  $\alpha$  from R ( $\alpha$  can specify multiple columns); it's distributive only over **union**
  - $R1 \times R2$  = select \* from R1 cross join R2 (cross product); the schema will contain all the atts of R1 followed by the atts of R2
  - $R1 \cup R2, R1 - R2, R1 \cap R2$  = union, set diff, intersection
  - $R1 \bowtie_\Theta R2$  = .. inner join on  $\Theta$  (condition join)
  - $R1 * R2$  = natural join; schema will contain the union of the atts (= all atts, **common ones appear only once**); returns 1 relation instance

- g.  $R1 \bowtie_c R2$  = left (outer) join on c; the reversed symbol is right (outer) join..
- h.  $R1 \bowtie_c R2$  = full (outer) join on c; schema will contain all the atts of R1 followed by the atts of R2
- i.  $R1 \triangleright R2$  = left semi join; the reverse is right semi join
- i. Left.. = the tuples in R1 that are used in the natural join; right.. = ..R2

j.  $R1 \div R2$  = division

k.  $\delta$  = duplicate elimination

l.  $S_{\{list\}}(R)$  = sort

m.  $\gamma_{\{list1\} group by \{list2\}}(R)$

n.  $\rho$  = renaming operator (nush ce plm e cu asta)

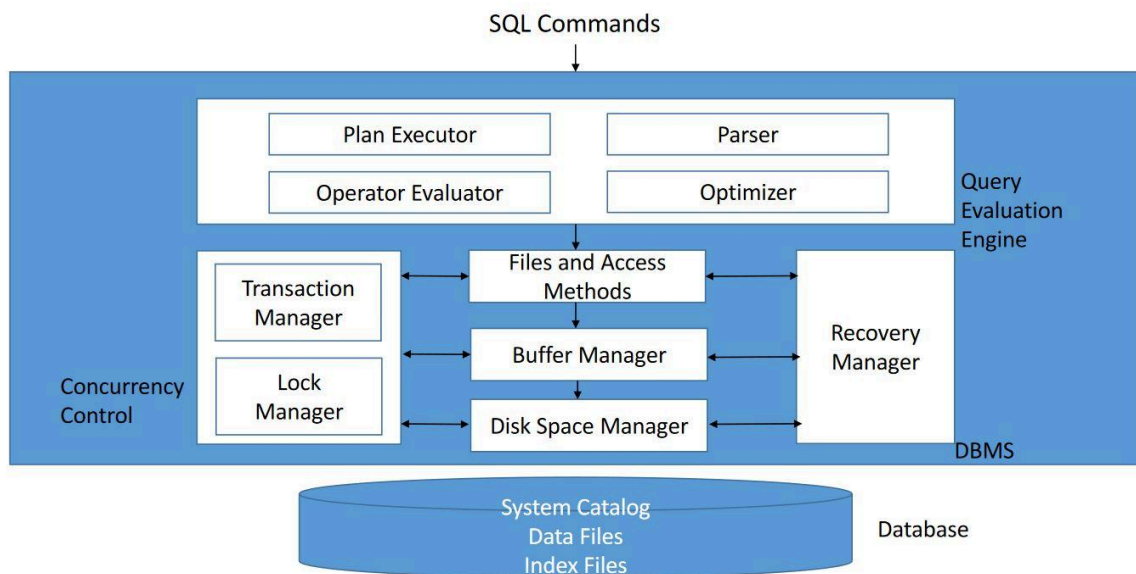
• for the previously described query language, with operators:  
 $\{\sigma, \pi, \times, \cup, -, \cap, \otimes, *, \bowtie, \bowtie_c, \triangleright, \triangleleft, \div\}$   
 an independent set of operators is  $\{\sigma, \pi, \times, \cup, -\}$

lecture 7

45. All the operators (without duplicate elim, sort, group by) can be obtained from  $\{\sigma, \pi, \times, \cup, -\}$  (= independent set of operators) (see 17. p3 on how to do it)

46. If there are multiple triggers defined for the same action, they are executed in a random order

47. DBMS structure



48. DBMSs operate on data when it is in memory

49. **Seek time fmm Carina** = time required to move the disk head to the desired track (smaller platter size => decreased seek time)

50. **Disk access time** (magnetic disk) = seek time + rotational delay + transfer time

51. **Block** = unit of data transfer between disk and main memory
52. **DSM** = disk space manager; monitors disk usage, it commands to (de)allocate, read / write a page
53. **BM = buffer manager**, brings new data pages from disk to main memory as they are required, manages the available main memory
  - a. **Replacement policies** - LRU (least recently used), MRU, random, clock replacement, toss-immediate
  - b. **Buffer pool (BP)** is made of frames, which can fit a page
  - c. Each frame has
    - i. **pin\_count** = no. of current users on the crt frame; only those with 0 can be replaced; initial value = 0
    - ii. **dirty** = true if the page has been changed since it was brought to the frame; initially off
  - d. When a page is requested
    - i. If the page is in the BP, increase its pin count => done
    - ii. Else, select a frame (FR) to replace using the replacement policies; if the old frame was dirty, write it on the disk; pin\_count(FR)++, the new page is read by the BM in that frame
    - iii. If the BP is full, the operation may be aborted, or it waits
  - e. It may pre-fetch pages
54. **Heap files**
  - a. Simplest file structure, unordered records, best when we need to scan all the records
  - b. Can use
    - i. Doubly-linked list - it needs a header = stored by the dbms; it has 2 lists, for full / free pages
    - ii. Directory of pages
55. **Index**
  - a. Stored on the disk, associated with a table or view
  - b. Speeds up equality / range select queries on the search key
  - c. If we change the data => the indexes need to be updated
  - d. hashed files - very good when searching for equality
  - e. **Data entry is:**
    - i. A1 - the actual data record with search key value = k
    - ii. A2 - <k, rid>, rid = id of a data record with sk = k
    - iii. A3 - <k, list of rid> ...
    - iv. In general, a2, a3 are smaller than a1
  - f. **Clustered**
    - i. A1 - always clustered

- ii. the order of the data records is close to / same as the order of the data entries
- iii. Includes all the columns in a table
- iv. At most 1 / collection of records
- v. When creating a PK, if there's no clus. on that table (and we don't specify unclustered) => create an unique cl. Index
- vi. Organized as a B+ tree

**g. Unclustered:**

- i. A2, A3 - are clus. only if the data records are ordered on the search key, which is uncommon in practice => A2, A3 are usually unclus
- h. **Primary index** - contains the PK; **unique index** = contains a candidate key
  - i. **secondary** = not primary (in our terminology, but in other places this might mean that it's just unclus); it can be unclus
- i. Primary / unique can't contain duplicates, secondary can
- j. Composite SK - contains several keys
- k. Covering index - contains all the columns that are necessary in a query
- l. Filtered index - can be created with a WHERE => only used when it satisfies that condition

**56. ISAM = Indexed Sequential Access Method**

- a. Not very good when doing lots of inserts / deletes
- b. Search complexity:  $\log_c m$  (c = children per index page, m = primary leaf pages)

**57. 2-3 tree**

- a. All terminal nodes on the same level
- b. If a non-terminal node has 1 key => 2 subtrees: one with values < key, the other > key
- c. If .. 2 keys (key1 < key2) => 3 subtrees: first < key1 < second < key2 < third


**d. Values in the tree are distinct**

**58. B-tree of order m**

- a. Generalization of 2-3 tree (= b-tree of order 3)
- b. All terminal nodes on the same level
- c. Every non-terminal, non-root node has between  $\text{ceil}(m / 2)$  and m **subtrees**
- d. A node with p subtrees has p-1 ordered **values**



- non-terminal, non-root node – at most 5, at least 3 subtrees (between 2 and 4 values)



59. **B+ tree**

- Variant of the B-tree
- Uniform search time; few I/O ops needed for searching; ideal for range selection
- Much better than ISAM for insert / delete etc.

60. **Hash-based indexing** - ideal for equality selections

61. **Static hashing** - the classic one (SDA gen)

62. **Extendible hashing** - there's a *directory* which points to the buckets, each determined by the last  $k$  (= *depth*) pro digits of the values; when a bucket becomes full, we split it into 2 buckets determined by the last  $k + 1$  digits and we double the size of the directory

63. The client generates SQL statements and sends them to the server, which syntactically analyses + evaluates them and sends the results back

64. **Table scan** - high transfer time for large tables (the slowest)

65. **Index seek** - when we search for a key value (using = ) that we have indexed (the fastest)

66. **Index scan** - the search condition can depend on either a key or a non-key att; some mem. blocks may be read multiple times

67. Joins are used in practice more often than cross-products. The sooner we apply selections / projections in the join implementation, the better

68. Some algorithms for **cross-join**, (indexed) nested loop join, merge join, hash join (found [here](#))

69. **Algebra equivalences**

a.  $\sigma_C(\pi_\alpha(R)) = \pi_\alpha(\sigma_C(R))$  (second is more efficient)

b.  $\sigma_{C1}(\sigma_{C2}(R)) = \sigma_{C1 \text{ AND } C2}(R)$  (second - more eff because it only does one pass of R)



- c.  $\sigma_C(R \times S) = R \otimes_C S$  (condition join - more eff)
- $\sigma_C(R \cup S) = \sigma_C(R) \cup \sigma_C(S)$
- $\sigma_C(R \cap S) = \sigma_C(R) \cap \sigma_C(S)$
- $\sigma_C(R - S) = \sigma_C(R) - \sigma_C(S)$
- d. (if the schemas of R and S are compatible)
- e.  $\sigma_C(R \times S) = \sigma_C(R) \times S$  (if C contains only atts from R)
- f.  $\sigma_{C1 \text{ AND } C2}(R \times S) = \sigma_{C1}(R) \times \sigma_{C2}(S)$  (if C1 contains only atts from R, C2 .. S)
- g.  $\sigma_{C1 \text{ AND } C2}(R \times S) = \sigma_{C1}(R \otimes_{C2} S)$  (if C2 is a join condition between R and S)
- h.  $\pi_\alpha(R \cup S) = \pi_\alpha(R) \cup \pi_\alpha(S)$
- i.  $\pi_\alpha(R \otimes_C S) = \pi_\alpha(\pi_{\alpha1}(R) \otimes_C \pi_{\alpha2}(S))$  ( $\alpha1$  = atts from R that appear in  $\alpha$  or C,  $\alpha2$  = ... from S...)

70. Intersection and union are associative and commutative
71. Cross-product and natural join are associative
72. When using the cross-join algorithm, the order of the data sources is important
73. **Surrogate key** = key that isn't obtained from the domain of the modeled problem
74. **DB design stages**
- Requirement analysis
  - Conceptual design
  - Logical design
  - Schema refinement (normalization, eliminate redundancy)
  - Physical design (ex. indexes)
75. We can have tables that reference themselves (recursive relationship); if there we have **ON DELETE CASCADE => error**; same if we have two tables that reference each other
76. The M:M table is also called *join table*
77. Reduce fragmentation
- In a heap - create then drop a clus index

- b. In an index - ALTER INDEX REORGANIZE (if small fragmentation), ALTER INDEX REBUILD (if large fragmentation) or drop + recreate index

78. **Indexed views**

- a. cannot reference other views
- b. the index must be clus + unique;
- c. only allowed aggregation operators are sum, count
- d. subqueries, outer joins, set operations etc are not allowed

SET options	required value	default server value
ANSI_NULLS	ON	ON
ANSI_PADDING	ON	ON
ANSI_WARNINGS	ON	ON
ARITHABORT	ON	ON
CONCAT_NULL_YIELDS_NULL	ON	ON
NUMERIC_ROUNDABORT	OFF	OFF
QUOTED_IDENTIFIER	ON	ON

e.

- 79. WAITFOR TIME '4:37' => execution continues at 4:37
- 80. Traditional dbms -> one-shot query (executed on the current instance of the data)
  - a. Human active, DBMS passive model (HADP)
- 81. **Data stream** = temporal sequence of values produced by a data source, potentially infinite; data is associated with timestamps; we can have data **stream** management sys
- 82. Event = elementary unit of information on a data stream (the equivalent of a record)
- 83. Sliding window = contiguous portion of the data stream, it has a size and a step size
- 84. **Continuous query** - perpetually running, DAHP